

laa9ynoje

November 10, 2025

```
[83]: import pandas as pd
import numpy as np
import plotly.express as px
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
from sklearn.metrics import confusion_matrix, classification_report, roc_curve, auc
import plotly.graph_objects as go
from plotly.subplots import make_subplots
```

```
[84]: data = pd.read_csv('credit_risk_data-1.csv')

data.head()
```

```
[84]:   application_id application_date  loan_amount  annual_income \
0          APP_2328      2022-01-01    132221.82    60451.82
1          APP_558       2022-01-01    134906.42   114634.08
2          APP_2477      2022-01-01     30285.19    82772.53
3          APP_741       2022-01-01     32516.09    94023.36
4          APP_145       2022-01-02     77900.99    53515.02

   employment_years  job_stability_score  credit_score  credit_utilization \
0                  6.6            0.898        679         0.106
1                 10.3            0.808        718         0.030
2                 12.1            0.964        768         0.174
3                  9.1            0.690        670         0.141
4                  7.2            0.679        651         0.097

   payment_history_score  open_credit_lines  debt_to_income_ratio \
0                0.876             1            0.451
1                0.719             4            0.090
2                0.775             6            0.201
3                0.993             3            0.322
4                0.946             2            0.222
```

```

    savings_ratio  asset_value  age education_level marital_status \
0          0.500     352569.55   41      High School       Married
1          0.235     224364.21   46        Masters      Divorced
2          0.172     514765.55   44      High School      Widowed
3          0.368     182541.72   26    Bachelors      Single
4          0.324     223691.29   50   Associates      Single

  residential_stability  loan_status
0                  3.5          0
1                 11.4          0
2                  8.6          0
3                  3.9          0
4                  9.6          0

```

[85]: `data.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2500 entries, 0 to 2499
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   application_id    2500 non-null   object  
 1   application_date   2500 non-null   object  
 2   loan_amount        2500 non-null   float64 
 3   annual_income      2500 non-null   float64 
 4   employment_years   2500 non-null   float64 
 5   job_stability_score 2500 non-null   float64 
 6   credit_score       2500 non-null   int64   
 7   credit_utilization 2500 non-null   float64 
 8   payment_history_score 2500 non-null   float64 
 9   open_credit_lines   2500 non-null   int64   
 10  debt_to_income_ratio 2500 non-null   float64 
 11  savings_ratio      2500 non-null   float64 
 12  asset_value        2500 non-null   float64 
 13  age                2500 non-null   int64   
 14  education_level    2500 non-null   object  
 15  marital_status     2500 non-null   object  
 16  residential_stability 2500 non-null   float64 
 17  loan_status         2500 non-null   int64   

dtypes: float64(10), int64(4), object(4)
memory usage: 351.7+ KB

```

[86]: `data.describe()`

```

[86]:      loan_amount  annual_income  employment_years  job_stability_score \
count    2500.000000     2500.000000     2500.000000     2500.000000
mean    155716.305344    67707.807596      6.675640      0.634643

```

```

std      149605.357952    27302.931731        3.488021        0.293276
min      5000.000000    15000.000000        0.000000        0.011000
25%     42984.517500    47475.317500        4.000000        0.375500
50%     97054.315000    66963.475000        6.700000        0.752000
75%    213214.992500    87347.642500        9.300000        0.866000
max     500000.000000   149929.960000       19.300000        0.999000

          credit_score  credit_utilization  payment_history_score \
count      2500.000000            2500.000000            2500.000000
mean      681.728400             0.358176            0.740733
std       88.683309             0.289995            0.285966
min      334.000000             0.004000            0.029000
25%     642.750000             0.131000            0.517500
50%     700.000000             0.246000            0.880500
75%     743.000000             0.592250            0.956000
max     850.000000             0.998000            1.000000

          open_credit_lines  debt_to_income_ratio  savings_ratio  asset_value \
count      2500.000000            2500.000000            2500.000000            2500.000000
mean      3.451600              0.408094            0.320784        175666.741236
std       2.083793              0.224736            0.192079        182652.568930
min      0.000000              0.009000            0.000000            550.630000
25%     2.000000              0.228000            0.161000        49513.082500
50%     3.000000              0.359000            0.327000        121018.750000
75%     5.000000              0.565000            0.464000        235513.902500
max     11.000000             0.979000            0.893000        1000000.000000

          age  residential_stability  loan_status
count      2500.000000            2500.000000            2500.000000
mean      42.045600             6.023200            0.265600
std       12.092395             3.205397            0.441741
min      18.000000             0.000000            0.000000
25%     34.000000             3.600000            0.000000
50%     42.000000             5.900000            0.000000
75%     50.000000             8.400000            1.000000
max     75.000000             16.400000            1.000000

```

[87]: `data.isna().sum()`

```

[87]: application_id      0
application_date      0
loan_amount           0
annual_income         0
employment_years      0
job_stability_score  0
credit_score          0
credit_utilization    0

```

```

payment_history_score      0
open_credit_lines          0
debt_to_income_ratio       0
savings_ratio              0
asset_value                0
age                         0
education_level             0
marital_status              0
residential_stability       0
loan_status                  0
dtype: int64

```

This database contains 17 columns, there is no null values in the database. In general this is a good database to work, as it contains several variables to work and clean data to determine loan_status.

```
[88]: fig = px.histogram(
        data,
        x='loan_status',
        title='<b>Distribution of Loan Status</b>',
        color='loan_status' # Optional: This gives each bar a different color
    )
fig.update_layout(
    title_x = 0.5
)

fig.show()
```

```
[89]: default_rate = data['loan_status'].mean()
print(f"The exact default rate is: {default_rate}")
```

The exact default rate is: 0.2656

```
[90]: key_predictors = ['credit_score', 'annual_income', 'debt_to_income_ratio',
                     'payment_history_score', 'savings_ratio']

#Visualization improved with the use of gemini
for col in key_predictors:
    # No inner loop needed.
    # Pass the full 'data' and use 'color' to differentiate.
    fig = px.histogram(
        data,
        x=col,
        color='loan_status', # Automatically creates colors for 0 and 1
        barmode='overlay',    # Layers the histograms (use 'group' for
                           # side-by-side)
        title=f'<b>Distribution of {col} by Loan Status</b>'
```

```

)
# Add opacity so the overlaid histograms are visible
fig.update_traces(opacity=0.75)

fig.update_layout(
    title_x=0.5,
    xaxis_title=col,
    yaxis_title='Count'
)
fig.show()

```

It is a clear difference in the distributions for the two loan_status types, this means both groups are really apart from each other.

```
[91]: education_levels = data['education_level'].unique()
marital_stats = data['marital_status'].unique()

[92]: education_defaults = data.groupby('education_level')['loan_status'].mean().
      ↪reset_index()

education_defaults = education_defaults.rename(columns={'loan_status': 'Mean\u20acDefault Rate'})

fig_edu = px.bar(
    education_defaults,
    x='education_level',
    y='Mean Default Rate',
    title='<b>Mean Default Rate by Education Level</b>',
    labels={'education_level': 'Education Level', 'Mean Default Rate': 'Average\u20acDefault Rate'}
)
# Sort the bars by the default rate (highest to lowest) for easier reading
fig_edu.update_layout(xaxis={'categoryorder': 'total descending'}, title_x = 0.5)

fig_edu.show()

[93]: marital_defaults = data.groupby('marital_status')['loan_status'].mean().
      ↪reset_index()

marital_defaults = marital_defaults.rename(columns={'loan_status': 'Mean\u20acDefault Rate'})

# Create the bar plot
fig_marital = px.bar(

```

```

    marital_defaults,
    x='marital_status',
    y='Mean Default Rate',
    title='<b>Mean Default Rate by Marital Status</b>',
    labels={'marital_status': 'Marital Status', 'Mean Default Rate': 'Average  
Default Rate'}
)
# Sort the bars by the default rate
fig_marital.update_layout(xaxis={'categoryorder':'total descending'}, title_x =  
0.5)

fig_marital.show()

```

```
[94]: all_num_predictors = ['loan_amount', 'annual_income', 'employment_years',  
                         'job_stability_score', 'credit_score', 'credit_utilization',  
                         'payment_history_score', 'open_credit_lines', 'debt_to_income_ratio',  
                         'savings_ratio', 'asset_value', 'age', 'residential_stability']

corr_mat = data[all_num_predictors].corr()

# heatmap con Plotly
fig = px.imshow(
    corr_mat,
    x=corr_mat.columns,
    y=corr_mat.index,
    color_continuous_scale='RdBu_r',
    zmin=-1, zmax=1,
    aspect="auto",
    text_auto=".2f"
)

fig.update_layout(
    title='<b>Correlation Matrix for all numeric predictors</b>',
    width=800,
    height=800,
    xaxis_title="Variables",
    yaxis_title="Variables",
    title_x =0.5
)

fig.show()
```

There are few variables with high collinearity, the one that have the most is payment_history_score with job_stability_score, this is the only correlation greater than 0.8.

```
[95]: #Data Preprocessing
print("Original data shape:", data.shape)
print("Original categorical variables:")
print("Education levels:", data['education_level'].unique())
print("Marital status:", data['marital_status'].unique())
```

Original data shape: (2500, 18)
 Original categorical variables:
 Education levels: ['High School' 'Masters' 'Bachelors' 'Associates' 'Doctorate']
 Marital status: ['Married' 'Divorced' 'Widowed' 'Single']

```
[96]: # Create dummy variables for categorical features
data_processed = pd.get_dummies(data, columns=['education_level', ↴
    'marital_status'], drop_first=True)

print("New data shape:", data_processed.shape)
print("Columns created:")
new_categorical_columns = [col for col in data_processed.columns if ↴
    'education_level_' in col or 'marital_status_' in col]
for col in new_categorical_columns:
    print(f" - {col}")
```

New data shape: (2500, 23)
 Columns created:
 - education_level_Bachelors
 - education_level_Doctorate
 - education_level_High School
 - education_level_Masters
 - marital_status_Married
 - marital_status_Single
 - marital_status_Widowed

```
[97]: # Drop non-predictive columns
x = data_processed.drop(['application_id', 'loan_status'], axis=1)
y = data_processed['loan_status']

print(f"Predictors x: {x.shape}")
print(f"Target y: {y.shape}")
print(f"N of features: {x.shape[1]}")
print(f"Feature names: {list(x.columns)}")
```

Predictors x: (2500, 21)
 Target y: (2500,)
 N of features: 21
 Feature names: ['application_date', 'loan_amount', 'annual_income',
 'employment_years', 'job_stability_score', 'credit_score', 'credit_utilization',
 'payment_history_score', 'open_credit_lines', 'debt_to_income_ratio',
 'savings_ratio', 'asset_value', 'age', 'residential_stability',

```
'education_level_Bachelors', 'education_level_Document', 'education_level_High School', 'education_level_Masters', 'marital_status_Married',  
'marital_status_Single', 'marital_status_Widowed']
```

```
[98]: #Train test split  
X_train, X_test, y_train, y_test = train_test_split(  
    x, y,  
    test_size=0.2,  
    random_state=42,  
    stratify=y  
)  
  
print(f"Training set X_train: {X_train.shape}, y_train: {y_train.shape}")  
print(f"Test set X_test: {X_test.shape}, y_test: {y_test.shape}")  
  
print(f"Target distribution in dataset:")  
print(y.value_counts(normalize=True))  
print(f"Target distribution in training set:")  
print(y_train.value_counts(normalize=True))  
print(f"Target distribution in test set:")  
print(y_test.value_counts(normalize=True))
```

```
Training set X_train: (2000, 21), y_train: (2000,)  
Test set X_test: (500, 21), y_test: (500,)  
Target distribution in dataset:  
loan_status  
0      0.7344  
1      0.2656  
Name: proportion, dtype: float64  
Target distribution in training set:  
loan_status  
0      0.7345  
1      0.2655  
Name: proportion, dtype: float64  
Target distribution in test set:  
loan_status  
0      0.734  
1      0.266  
Name: proportion, dtype: float64
```

```
[99]: # Clean date column and standardize  
if 'application_date' in X_train.columns:  
    X_train['application_date'] = pd.to_datetime(X_train['application_date'])  
    X_test['application_date'] = pd.to_datetime(X_test['application_date'])  
    ref_date = X_train['application_date'].max()  
    X_train['app_date_days'] = (ref_date - X_train['application_date']).dt.days  
    X_test['app_date_days'] = (ref_date - X_test['application_date']).dt.days
```

```

X_train = X_train.drop('application_date', axis=1)
X_test = X_test.drop('application_date', axis=1)

# Standardization
scaler = StandardScaler()
X_train_scaled = pd.DataFrame(scaler.fit_transform(X_train), columns=X_train.
                                columns, index=X_train.index)
X_test_scaled = pd.DataFrame(scaler.transform(X_test), columns=X_train.columns,
                                index=X_test.index)

print(f"Data standardized Train: {X_train_scaled.shape}, Test: {X_test_scaled.
                                shape}")

```

Data standardized Train: (2000, 21), Test: (500, 21)

0.1 Statistical Assumption Testing

Statistical Assumptions for LDA vs. QDA

Linear Discriminant Analysis and Quadratic Discriminant Analysis are based on the assumption of multivariate normality that predictors follow approximately a multivariate normal distribution within each class (defaulters vs. non-defaulters).

Evaluation of Assumptions Based on EDA:

Multivariate Normality: From our EDA histograms, we observe that: - credit_score: Shows approximately normal distribution for both classes - annual_income: Reasonably normal distribution with some right skew - debt_to_income_ratio: Approximately normal for non-defaulters, slightly skewed for defaulters

While not all variables are perfectly normal, both LDA and QDA are reasonably robust to moderate deviations from normality, especially with our sample size of 2,500 observations.

Homogeneity of Covariance Matrices (Key Differentiating Assumption): - LDA assumes that all classes share the same covariance matrix - QDA relaxes this assumption and allows each class to have its own covariance matrix

Evidence from our EDA: - The distribution plots show clear differences in variance between defaulters and non-defaulters for several variables - For example, in credit_score, defaulters show wider spread and lower mean compared to non-defaulters - Similar pattern observed in annual_income and other financial metrics

Hypothesis: Given the visible differences in distributions and variances between the two classes in our EDA plots, we expect that the covariance matrices are unequal. Therefore, we hypothesize that QDA will outperform LDA as it can better capture these class-specific covariance structures.

```

[100]: lda = LinearDiscriminantAnalysis()
lda.fit(X_train_scaled, y_train)

# Check model performance on training data
train_accuracy = lda.score(X_train_scaled, y_train)

```

```

test_accuracy = lda.score(X_test_scaled, y_test)

print(f'LDA Training Accuracy: {train_accuracy:.4f}')
print(f'LDA Test Accuracy: {test_accuracy:.4f}')

#Interpret Coefficients
lda_coef = pd.DataFrame({
    'Feature': X_train_scaled.columns,
    'Coefficient': lda.coef_[0]
})

#Sort by absolute value to identify most important features
lda_coef['Abs_Coefficient'] = np.abs(lda_coef['Coefficient'])
lda_coef_sorted = lda_coef.sort_values('Abs_Coefficient', ascending=False)

print("Top 10 Most Important Features from LDA:")
print(lda_coef_sorted.head(10).to_string(index=False))

```

LDA Training Accuracy: 1.0000

LDA Test Accuracy: 1.0000

Top 10 Most Important Features from LDA:

	Feature	Coefficient	Abs_Coefficient
payment_history_score	-15.471241	15.471241	
job_stability_score	-13.051790	13.051790	
credit_utilization	11.766475	11.766475	
debt_to_income_ratio	4.471469	4.471469	
credit_score	-3.980499	3.980499	
savings_ratio	-2.995611	2.995611	
employment_years	-2.367811	2.367811	
residential_stability	-1.702223	1.702223	
annual_income	-1.587700	1.587700	
open_credit_lines	-1.275071	1.275071	

[101]: # Create a bar plot of top coefficients

```

fig_lda_coef = px.bar(
    lda_coef_sorted.head(10),
    x='Coefficient',
    y='Feature',
    orientation='h',
    title='<b>Top 10 Most Important LDA Coefficients</b>',
    color='Coefficient',
    color_continuous_scale='RdBu_r',
    labels={'Coefficient': 'Standardized Coefficient Value', 'Feature': ''})
)
fig_lda_coef.update_layout(
    title_x=0.5,
    yaxis={'categoryorder': 'total ascending'},

```

```

        xaxis_title="Coefficient Value",
        yaxis_title="Feature",
        showlegend=False
    )

fig_lda_coef.show()

```

0.2 LDA Coefficients Interpretation

The coefficients tell us the importance and direction of each variable's relationship with default risk, a big negative coefficient reduces default risk while a big positive one increases it.

Key Drivers of Default Risk:

1. **payment_history_score (Coefficient: -15.47)**: This is, by a wide margin, the most important predictor the negative sign indicates that a higher payment history score is strongly associated with a lower probability of default (being a good payer).
2. **job_stability_score (Coefficient: -13.05)**: The second most important driver, similarly, a higher job stability score dramatically decreases the likelihood of default.
3. **credit_utilization (Coefficient: 11.77)**: This is the strongest positive driver, the positive sign means that higher credit utilization (using a higher percentage of your available credit) is strongly associated with higher default risk.
4. **debt_to_income_ratio (Coefficient: 4.47)**: A positive coefficient indicating that as an applicant's debt to income ratio increases so does their probability of default.
5. **credit_score (Coefficient: -3.98)**: As expected, a higher credit score has a negative sign, reducing the likelihood of default.

Key Insights: The model identifies a clear risk profile. Default risk increases massively for applicants who use a large amount of their available credit and carry a high debt burden relative to their income. Conversely, risk drops significantly for those with a strong payment history and stable employment.

```
[102]: qda = QuadraticDiscriminantAnalysis()
qda.fit(X_train_scaled, y_train)

# Check model performance
qda_train_accuracy = qda.score(X_train_scaled, y_train)
print(f"QDA Training Accuracy: {qda_train_accuracy:.4f}")
qda_test_accuracy = qda.score(X_test_scaled, y_test)
print(f"QDA Test Accuracy: {qda_test_accuracy:.4f}")
```

QDA Training Accuracy: 0.9995

QDA Test Accuracy: 1.0000

```
[103]: #predictions and probability scores
y_pred_lda = lda.predict(X_test_scaled)
y_pred_qda = qda.predict(X_test_scaled)
y_score_lda = lda.predict_proba(X_test_scaled)[:, 1]
y_score_qda = qda.predict_proba(X_test_scaled)[:, 1]
```

```
[104]: #subplots for confusion matrices
fig = make_subplots(
    rows=1, cols=2,
    subplot_titles=['LDA Confusion Matrix', 'QDA Confusion Matrix'],
    horizontal_spacing=0.15
)

# LDA
cm_lda = confusion_matrix(y_test, y_pred_lda)
heatmap_lda = go.Heatmap(
    z=cm_lda,
    x=['Predicted 0', 'Predicted 1'],
    y=['Actual 0', 'Actual 1'],
    text=cm_lda,
    texttemplate="%{text}",
    textfont={"size": 16},
    colorscale='Blues',
    showscale=False
)
fig.add_trace(heatmap_lda, 1, 1)

# QDA
cm_qda = confusion_matrix(y_test, y_pred_qda)
heatmap_qda = go.Heatmap(
    z=cm_qda,
    x=['Predicted 0', 'Predicted 1'],
    y=['Actual 0', 'Actual 1'],
    text=cm_qda,
    texttemplate="%{text}",
    textfont={"size": 16},
    colorscale='Blues',
    showscale=False
)
fig.add_trace(heatmap_qda, 1, 2)

fig.update_layout(
    title_text="Confusion Matrices: LDA vs QDA",
    title_x=0.5,
    width=800,
    height=400
)

fig.show()
```

```
[105]: print("LDA Classification Report")
print(classification_report(y_test, y_pred_lda))
```

```

print("QDA Classification Report")
print(classification_report(y_test, y_pred_qda))

```

LDA Classification Report				
	precision	recall	f1-score	support
0	1.00	1.00	1.00	367
1	1.00	1.00	1.00	133
accuracy			1.00	500
macro avg	1.00	1.00	1.00	500
weighted avg	1.00	1.00	1.00	500

QDA Classification Report				
	precision	recall	f1-score	support
0	1.00	1.00	1.00	367
1	1.00	1.00	1.00	133
accuracy			1.00	500
macro avg	1.00	1.00	1.00	500
weighted avg	1.00	1.00	1.00	500

```

[106]: #Calculate ROC curves and AUC
fpr_lda, tpr_lda, _ = roc_curve(y_test, y_score_lda)
fpr_qda, tpr_qda, _ = roc_curve(y_test, y_score_qda)
auc_lda = auc(fpr_lda, tpr_lda)
auc_qda = auc(fpr_qda, tpr_qda)

print(f'LDA AUC: {auc_lda:.4f}')
print(f'QDA AUC: {auc_qda:.4f}')

fig_roc = go.Figure()
# LDA ROC curve
fig_roc.add_trace(go.Scatter(
    x=fpr_lda, y=tpr_lda,
    mode='lines',
    name=f'LDA (AUC = {auc_lda:.4f})',
    line=dict(width=3, color='blue')
))

# QDA ROC curve
fig_roc.add_trace(go.Scatter(
    x=fpr_qda, y=tpr_qda,
    mode='lines',
    name=f'QDA (AUC = {auc_qda:.4f})',

```

```

        line=dict(width=3, color='red')
    )))
#random classifier line
fig_roc.add_trace(go.Scatter(
    x=[0, 1], y=[0, 1],
    mode='lines',
    name='Random Classifier',
    line=dict(dash='dash', color='black')
))
fig_roc.update_layout(
    title='<b>ROC Curves: LDA vs QDA</b>',
    title_x=0.5,
    xaxis_title='False Positive Rate',
    yaxis_title='True Positive Rate',
    width=700,
    height=500,
    showlegend=True
)
fig_roc.update_xaxes(range=[0, 1])
fig_roc.update_yaxes(range=[0, 1])
fig_roc.show()

```

LDA AUC: 1.0000

QDA AUC: 1.0000

In this analysis, we compared two classification models, Linear Discriminant Analysis and Quadratic Discriminant Analysis, to predict loan default risk, the evaluation was performed on a held out test set (20% of the data).

Test Set Performance Summary:

Metric	LDA Model	QDA Model
Accuracy	1.0000	1.0000
AUC	1.0000	1.0000
Recall (Class 1 - Default)	1.00	1.00
Precision (Class 1 - Default)	1.00	1.00

Results Analysis: Both models achieved perfect performance on the test set, with 100% accuracy and an AUC of 1.0000, the confusion matrices and classification reports confirm that both models correctly classified every single defaulter and non defaulter in the test set.

This result indicates the dataset is perfectly separable. As noted in the EDA and confirmed by the LDA coefficients, there are features (most notably payment_history_score and job_stability_score) that draw a perfect line between the two classes.

Model Selection: Given that both models have identical, perfect performance, either could be chosen.

The LDA Model is preferable for its interpretability, as it provides direct coefficients that explain why an applicant is flagged as high risk, which is invaluable for the business. QDA Model confirmed our hypothesis from Section 4 (that the covariance matrices were unequal), but its complexity (a quadratic decision boundary) is unnecessary when a simple linear boundary (LDA) already solves the problem perfectly.

Technical Recommendation: I select the LDA model, it achieves perfect accuracy while offering the benefit of model interpretability, allowing LendSmart to understand the key factors driving risk.